

Project 1: Cryptographic Voting

This project has two due dates: **Friday, February 19** and **Friday, February 26**, both at **6 p.m.**. Late submissions are not accepted. If you have a conflict due to travel, interviews, etc., please plan accordingly and turn in your homework early.

This is a group project; you will work in **teams of two** and submit one project per team. Please find a partner as soon as possible. If have trouble forming a team, post to the relevant Piazza thread. The final exam will cover project material, so you and your partner should collaborate on each part.

The code and other answers your group submits must be entirely the work of your group, and you are bound by the Honor Code. You may consult with other students about the conceptualization of the project and the meaning of the questions, but you may not look at any part of someone else's solution or collaborate with anyone outside your group. You may consult published references, provided that you appropriately cite them (e.g., with program comments), as you would in an academic paper. *Even though this project is based on code subject to an open source license, you may not redistribute your work, because that would constitute a violation of the Honor Code.*

Once you have your partner, then you will use GitHub Classroom, via this clone link. One of you goes first, and GitHub will ask for a name for your “team”. Please use your two Rice NetIDs with a hyphen between them, e.g., abc12-def34, as your team name. **Do not visit the clone link until both you and your partner have agreed to work together.** GitHub makes it difficult to rearrange partnered repositories, so we'd like to avoid this unless absolutely necessary.

Before you do anything else, be sure to edit the README.md file to have your names, email addresses, and so forth, then commit the file. Just to make sure you're not having any weird GitHub issues, do this with separate commits: one from each partner editing their name. That way, if there are weird GitHub problems, we'll know about them as soon as possible.

<https://classroom.github.com/a/XXXXXXXXX>

1 Introduction

This project gives you a sophisticated toolbox of primitives that are useful for building “end to end” cryptographic voting systems, including ElGamal cryptography (additively homomorphic) and Chaum-Pedersen proofs. It's based on an open-source project from Microsoft called ElectionGuard,

which has multiple implementations for different purposes. This assignment is based on the Python “reference implementation”¹ but is implemented in Kotlin.

The “real” ElectionGuard code has a number of features that are useful for “real” elections, like understanding different “ballot styles” (i.e., not every candidate appears on every ballot, but you still want to be able to add ballots together). These features add a huge amount of complexity that might get in your way of understanding and using the code without a lot of study.

Consequently, we’ve stripped all the fancy stuff out and kept all the cryptographic primitives. You’ll be implementing “simple” elections, where every ballot has exactly one contest on it, with the same candidates on every ballot, and each voter gets to pick k of the n total candidates. These are simply added up, and whoever gets the most wins. (This is a variation on “approval voting”², which is something of a generalization of the usual “vote for one” scheme, sometimes called “first past the post”³.)

This project has two parts. You’ll first encrypt, decrypt, and validate “selections”, which is to say, the individual choices that a voter might make. For the second part, you’ll move up one level of abstraction and implement encryption, decryption, and validation for whole ballots, as well as for the election tabulations.

For both parts, we’ve written extensive unit tests, leveraging the Java QuickTheories property-based testing library⁴, which will exercise your code with both positive examples (i.e., where validation should succeed) and negative examples (i.e., where validation should fail). We’re also going to ask you to write some of your own unit tests that validate specific properties. Your grade will be primarily based on which unit tests you pass and which unit tests you fail. Of course, you’re welcome to add any additional tests you want, if you find that helpful, but *if you make any modifications to existing tests, we will grade them as if they failed.*

1.1 Objectives

- Gain exposure to modern, real-world cryptographic techniques.
- Learn to implement cryptographic code and validate its correctness.

¹<https://github.com/microsoft/electionguard-python>

²https://en.wikipedia.org/wiki/Approval_voting

³https://en.wikipedia.org/wiki/First-past-the-post_voting

⁴<https://github.com/quicktheories/QuickTheories>

2 Setting up your computer

This assignment generally requires you to set up your computer as one would do for a modern Kotlin development environment, which pretty much means installing IntelliJ IDEA along with any OpenJDK. You'll also be able to run all the tests from the command-line, and that's how we'll describe everything here.

Once your computer is ready, you should install:

- Any OpenJDK 8 or later. Azul notably provides builds that work on ARM computers like M1 Macs. (<https://www.azul.com/downloads/?package=jdk>)
- (Mac only) Install the Apple developer tools. Run `xcode-select -install` and click to accept the license.
- (Windows only) Git (<https://gitforwindows.org/>) – this includes the “Git Bash Shell”, where you can type the commands we'll show you.
- (Windows only) A fancier way to install Git is through Chocolatey, which also lets you easily install other commands you might want (Simple instructions: <https://jcutrer.com/windows/install-chocolatey-choco-windows10>, more complex instructions: <https://chocolatey.org/install>).

3 Cloning from GitHub, setting up your environment

If you're using a tool like IntelliJ, it will do all of this automatically. You give it the GitHub URL, and it will clone it locally. You may need to tell IntelliJ where you installed your JDK, otherwise everything should “just work”.

Once installed, you should run the unit tests, either from IntelliJ, or by running `./gradlew test` on the command-line. Most of the unit tests should succeed, with others that will fail because you haven't implemented the assignment yet.

Similarly, you can compare what you see on your own computer with what happens every time you push some commits to GitHub, where all the unit tests will be executed in a virtual environment by GitHub Actions. You'll see a red X or green checkmark, telling you whether the tests have failed or passed. You can click on that and you'll then see a transcript of the unit tests running.

4 Cryptographic basics

There are several cryptographic things going on in ElectionGuard that we didn't have time to cover in class, but that you really need to understand.

When we introduced the digital signature algorithm, we mentioned that it has two separate primes that it uses, p and q , but we didn't spend much time explaining why, except to say that we

want our exponents to be smaller, if we can get away with it without giving up any cryptographic strength. (The text below was written by Josh Benaloh, a principal senior cryptographer at Microsoft Research, who also designed the cryptography for ElectionGuard.)

The reason for two primes is that if you are doing computations in the multiplicative subgroup modulo p , the elements are $1, 2, \dots, p-1$. The identity 1 is always special, but since $p-1$ will not be prime (unless $p=3$), some of the remaining elements will be able to be *distinguished*. All elements x will have the property $x^{p-1} \bmod p = 1$, but many have smaller order as well. For instance, every $x = y^2 \bmod p$ (quadratic residues) has the property that $x^{(p-1)/2} \bmod p = 1$, and this property persists as these elements are raised to other powers. This can compromise information that should be kept secret.

To avoid having distinguished elements, we like to use a *prime order subgroup*. In many cases, we choose p such that $(p-1)/2$ is prime and just use the prime order subgroup consisting of all quadratic residues. However, we can achieve the security of a very large p by using a much smaller subgroup defined by a smaller q that divides $p-1$. All of the elements of this smaller subgroup are generated by $g^0, g^1, g^2, \dots, g^{q-1}$, so we can get random elements with size q exponents rather than size p exponents. This makes everything substantially more efficient. A q just below 2^{256} also allows us to match exponents nicely with SHA-256 outputs. So using q as the largest 256-bit prime works very nicely.

One might then choose p to be the largest 4096-bit prime such that $p-1$ is a multiple of q (which is what I did initially). However, having p be close to a power of 2 (or otherwise representable as a value of a polynomial with “small” coefficients) weakens its resistance to discrete log attacks via the special number-field sieve⁵. It is desirable to have parameters generated deterministically to avoid concerns that there may be some hidden back door. I tried to find a clean way to do this, but every attempt that I made wound up being representable as a polynomial with small coefficients. I therefore did what everyone else does and used an unrelated mathematical constant to deterministically produce p . Both π and e are overused for this purpose, and the description is slightly cleaner if the constant is less than one and greater than one half, so I wound up with the Euler-Mascheroni constant⁶. For computational efficiency, there is some benefit to having all of the extreme high-order bits and low-order bits be one, so this drove the selection of a suitable p .

5 ElGamal encryption

Normal “exponential” ElGamal encryption, invented in 1985 by Tahir ElGamal⁷, is very straightforward. If you have a secret key a in $[0, Q)$, the corresponding public key is $g^a \bmod p$, in $[0, P)$. (We’ll stop writing the P , Q , and \bmod parts because they just overcomplicate the equations.) It’s safe to share g^a with anybody because it’s (hopefully) very hard to recover a . This is called the “discrete log problem” (see also, the not entirely identical “decisional Diffie-Hellman problem”⁸, which is the actual hard problem we’re counting on for ElGamal encryption).

⁵https://en.wikipedia.org/wiki/Special_number_field_sieve

⁶https://en.wikipedia.org/wiki/Euler-Mascheroni_constant

⁷<https://ieeexplore.ieee.org/document/1057074>

⁸https://en.wikipedia.org/wiki/Decisional_Diffie%E2%80%93Hellman_assumption

Because we're interested in having an *additive homomorphism*, where we can “add” together two encryptions and get another encryption, we're going to use a variant called “exponential ElGamal”. With this, we put the message m we're encrypting into an exponent. For voting, these messages are mostly going to be 0-or-1, representing a voter's selection or absence of a selection, so it's important that the cryptosystem have an input that ensures a wide distribution of possible ciphertexts. That's a random number r called a nonce. We can now define encryption and decryption as follows:

$$\begin{aligned}
c &= \text{Encrypt}(m, r, g^a) = \langle g^r, g^{ar} g^m \rangle \\
p &= \text{Decrypt}(c, a) \\
&= \text{Decrypt}(\langle g^r, g^{ar} g^m \rangle, a) && \text{note : } (g^r)^a = g^{ar} \\
&= \text{DLog}_g \left(g^{ar} g^m * (g^{ar})^{-1} \right) \\
&= \text{DLog}_g g^m \\
&= m
\end{aligned}$$

Note that $\text{DLog}_g g^m$ is only efficiently computable for relatively small m , since it requires building a table $[1, g, g^2, g^3, \dots]$. Multiplicative inverses, however, are always efficient to find, for any value in the group.

5.1 Homomorphic accumulation

With this definition of ElGamal encryption, we can define an accumulation operation using piecewise multiplication:

$$\begin{aligned}
\text{sum} &= \text{Encrypt}(m_1, r_1, g^a) \oplus \text{Encrypt}(m_2, r_2, g^a) \\
&= \langle g^{r_1}, g^{ar_1} g^{m_1} \rangle \oplus \langle g^{r_2}, g^{ar_2} g^{m_2} \rangle \\
&= \langle g^{r_1+r_2}, g^{a(r_1+r_2)} g^{m_1+m_2} \rangle \\
&= \text{Encrypt}(m_1 + m_2, r_1 + r_2, g^a)
\end{aligned}$$

Any intermediary can compute the homomorphic sum of these encryptions without needing to decrypt, and without needing to know the secret key.

5.2 Key sharing

The real ElectionGuard supports a sophisticated “threshold cryptography” system which allows multiple trustees to each have a “share” of the election's key material. These shares are combined into a single public key, that all the voting machines use. For decryption operations, it takes a (predefined) k of the n trustees to collaborate together in order to do any decryption operation. This “key ceremony” is a bit complicated.

As a warmup, you'll be implementing a simplified version, where all of the key shareholders need to work together. Here's how it works. Let's say we have n trustees, and each has secret keys a_i and public keys g^{a_i} . We can define a joint public key by simply multiplying them all together:

$$\begin{aligned} k_{\text{joint}} &= g^{a_0} g^{a_1} \dots g^{a_n} \\ &= g^{a_0 + a_1 + \dots + a_n} \end{aligned}$$

Multiplying together the public keys yields a joint key, where we know the *formula* for the exponent, even though we can't derive the *value* for it. So, if a message is encrypted with this joint key, we need a way for the trustees to collaboratively decrypt it.

$$\begin{aligned} \text{PartialDecrypt}(c, a_i) &= \text{PartialDecrypt}\left(\left\langle g^r, g^{(a_0 + a_1 + \dots + a_n)r} g^m \right\rangle, a_i\right) \\ \text{DecryptionShare}_i &= (g^r)^{a_i} = g^{a_i r} \\ \prod_i \text{DecryptionShare}_i &= g^{a_0 r} g^{a_1 r} \dots g^{a_n r} \\ &= g^{(a_0 + a_1 + \dots + a_n)r} \end{aligned}$$

Each decryption share is nothing more than the nonce padding g^r raised to the secret key share a_i , and then we can just multiply those shares together and we get exactly the value that's multiplied with g^m . As before, multiplicative inverses are easy to find, so we have everything we need to combine "partial decryptions" together to yield the complete decryption. (This is a warm-up exercise that you'll do to get things rolling for part 1 of the assignment.)

5.3 Chaum-Pedersen

You'll see several different kinds of Chaum-Pedersen proof objects in `ChaumPedersen.kt`:

DisjunctiveChaumPedersenProofKnownNonce demonstrates that the prover knows the secret nonce used to create an encryption that's either a 1 or 0.

ConstantChaumPedersenProofKnownNonce demonstrates that the prover knows the secret nonce used to create an encryption that's a constant (not just 0 or 1).

ConstantChaumPedersenProofKnownSecretKey demonstrates that the prover knows the secret key used to create an encryption that's a constant (not just 0 or 1).

GenericChaumPedersenProof is used by all of the above proofs. It provides a proof that two tuples $(g, g^x), (h, h^x)$ share the same exponent x without revealing it.

So, when do you use which sort of proof? When you are *encrypting a counter that might be one or zero*, you make a disjunctive proof. You'll need to hang on to the nonce that you used to do the encryption. You'll ultimately output one proof per encrypted counter.

When you are *proving that the voter cast zero or one votes on the ballot*, you could theoretically use a disjunctive proof again, but we want to support k -of- n candidates being selected. That would

make a disjunctive proof much more complicated. Instead, what you'll do is create zero or more "placeholder" candidates (literally, their name will be "PLACEHOLDER" with an integer after that). If all the other candidates got zero votes, then the placeholder(s) gets one. Otherwise, the placeholder(s) gets zero, such that the sum of the candidate slots plus the placeholder slots is exactly k . Each placeholder ciphertext has the same disjunctive proof as all the real candidates. And now here's the cool part: you just use the additive homomorphism of ElGamal, and generate a *constant* Chaum-Pedersen proof, that the sum of all the candidates, including the placeholder, is exactly k .

When you are *decrypting the tally and proving that the announced plaintext is correct*, you use the decryption proof, which doesn't require the nonce, but does require the decryption (secret) key. You'll see each of these Chaum-Pedersen proof types listed as part of the various data structures in `SimpleElectionData.kt`, so you'll know which one to use where.

So, how do those proofs work? We cover this in more detail in our lecture. For the purposes of this assignment, what matters is that you understand what a "generic" Chaum-Pedersen proof actually proves, which then enables the more specific versions of the proof to use it as a subroutine. So, if we want to prove that the encryption of $m = 1$ is actually one, based on our knowledge of the encryption nonce r , then we're going to be generating a proof that (g, g^r) and (g^a, g^{ar}) share the nonce r without revealing r . And that value g^{ar} is just the right-hand side of the ElGamal encryption, after we multiply it with g^{-1} (i.e., canceling out the g^m where $m = 1$).

This might not seem very exciting, but it's the same math when we're proving that the *sum of the encrypted counters* is exactly the constant k for a k -of- n election. When you combine that with proofs, for each counter, that they are exactly 0 or 1, you end up with a proof, publicly verifiable, that you have a *well-formed ballot*. No overvotes. No funny business. This gives public observers of an election the ability to verify important properties of the election tally without compromising any individual ballot's privacy!

5.4 Nonces and nondeterminism

You'll see `seed`, `hashHeader`, and other such names for arguments to many of your functions. The important concept is that **all your code should be deterministic**. The same inputs should always yield the same outputs. These seed values are where you get your nondeterminism. When you're encrypting a ballot that has five counters, you'll need five individual random numbers. You can derive them from the seed using the code in `Nonces.kt`.

The broad idea, then, is that every "real world" ballot will start with a single, true-random number, and from that, all of the other random r values for the encryptions, and everywhere else they're needed, will be derived from that initial seed.

You want to implement this as a hierarchical process. From the master seed, you'll derive seeds for each ballot. Then from the ballot seeds, you'll derive seeds for each selection. Each layer of code passes the seed values down to the next layer of code. You can convert one seed into as many as you need by using the code in `Nonces.kt`.

Extra fun. Imagine that we have hand-marked paper ballots being fed into a scanner with a reasonably fast computer but not enough storage for the ballot ciphertexts (real world ElectionGuard

ciphertexts can be multiple megabytes of data per ballot). If we printed a unique master seed as a barcode on the bottom of each ballot, then the scanner could compute the full encryption, if it wanted, printing the voter a “receipt” which is really just the hash of their encrypted ballot. *Then the machine can completely forget the ciphertext* because it can be recomputed later, identically, due to the deterministic architecture of how nonces are derived from seeds.

5.5 The ElectionGuard library

If you look in the `src/main/kotlin/electionguard` directory, you’ll see a number of files. Here’s what each of them does:

ChaumPedersen.kt Implementations of several different Chaum-Pedersen proofs (see Section 5.3).

DLog.kt Implements a caching, brute-force discrete log engine. So long as the exponent is on the small side (millions, not billions), this will run efficiently. We need this to undo the g^v exponentiation applied to the plaintext vote counter v .

ElGamal.kt Implements “exponential ElGamal” encryption and decryption. Also, this is where you’ll implement “key sharing” (see Section 5.2).

Group.kt Implements `ElementModQ`, `ElementModP`, and a variety of function wrappers to compute with them. All of this bottoms out to calls to the Java `BigInteger` class, which has fast, hand-tuned C and assembly code.

Hash.kt Implements a wrapper around SHA-256 hashing, where you can hash just about anything, including lists, and get out an `ElementModQ`.

Nonces.kt Often you’ll want to start with some seed value and generate a stream of nonces that you need to be different from one another, but deterministically derived from the seed. This code gives you a Kotlin class instance, which you can treat as if it’s an array of infinite length, giving you random access to any element. Optional “headers” (typically, just another string) will be hashed together with the given seed value, making it easy for you to create multiple nonce streams, which you might use for different purposes, from the same original seed value (see Section 5.4).

SimpleElectionData.kt This defines all the input and output types for the code you’ll be writing. This is the Kotlin equivalent of C struct definitions. (We tried to avoid using too many of Kotlin’s crazier functional or object-oriented features, to make our code more readable by non-experts, although we do use Kotlin extension functions extensively, so you may want to read up on that.)

SimpleElections.kt This is where you’ll be doing most of your work. We’ve defined all the functions you need to implement, including their argument and return types, but we’ve left most of the bodies empty.

Utilities.kt A handful of general-purpose utility functions that are used in various places.

But that's not all. You'll find corresponding unit test files in the `test` directory. If you read those tests, you'll see that they make extensive use of the QuickTheories property-based testing library, which allows inputs to be generated at random. In `src/test/kotlin/electionguard/Generators.kt`, you'll see the "generator" functions that can produce these random values. Unfortunately, QuickTheories normally only likes to generate at most four arguments at a time, but we often need more than that, so we had to cobble together "tuples" of various sizes for QuickTheories to generate more things at once. It make the code a little bit uglier, but it all works.

6 Submission Strategy & Engineering

All of the code elements that you must implement are conveniently labeled TODO for you and will mention “part 1” or “part 2”. You’ll notice some of these TODO comments are in the unit tests, themselves. Each TODO item explains what you’re supposed to do, so we won’t reiterate those here.

Generally speaking, part 1 is a warm-up, where you’re learning your way around the codebase, and part 2 is where you’re dealing with ballots having k-of-n votes cast. Of note, in part 1, you’ll start off by implementing ElGamal key sharing (see 5.2), which involves editing code in ElGamal.kt. Everything else you’ll do is in SimpleElections.kt and the associated unit test files.

You should adopt a “commit and push, early and often” strategy. While we’re not specifically instructing you in how you should collaborate with your partner, some variant of pair programming will be invaluable, where one of you is typing and sharing their screen, while the other is looking over and offering suggestions. Certainly, Git makes it straightforward for you to work on separate items and merge them later, but you’re less likely to understand the full project. Trade-off who is typing and who is assisting. Don’t just say “I know how to do this. You just watch and maybe clean it up later.”

The unit tests will all run on GitHub as part of GitHub Actions. It’s generally faster and easier to run them on your own computer. You can do this inside your favorite IDE, or you can run `./gradlew test` from your commandline. You should also run `./gradlew formatKotlin` to make sure your code is nicely formatted.

The autograder is also available to you from the commandline. Just run `./gradlew clean autograder`. That should get you results identical to what the graders see on GitHub Actions.

When you finish part 1. You’re going to want to use a Git tag. You do this after you’re satisfied that you’ve completed the requirements. You then say `git tag part1` and `git push -tags`. Alternately, you can use the GitHub website to create a “release”, which in turn creates a Git tag. Your grader will examine your code, at this tag, and the output of GitHub Actions when running the unit tests on this tag.

Continuing with part 2. You can march right ahead with part 2, even if that means breaking something from part 1. We will grade whatever you have committed and pushed prior to the deadline. If you want to keep playing with the code afterwards, please make a `part2` tag, as above, and then that will be the focus of the graders’ attention.

What if it works on my computer and fails on GitHub Actions? Your grade is based on what the grader sees running on GitHub Actions. Perhaps you changed a file on your local machine and didn’t commit or push the changes to the server. When in doubt, find one of the TAs or the professor and ask for help if you’re seeing these sorts of discrepancies.

Am I allowed to add imports to the files? Absolutely. If there’s a function you need that isn’t already imported, then feel free to import it. That includes functions from the ElectionGuard library,

or otherwise available from the standard Kotlin libraries. You shouldn't need to add or change any external dependencies in `build.gradle`.

I'm seeing a ton of warnings on the console... Many of the functions that you're given, when they discover an error, will log about it, and then indicate a failure as part of their return value. Some unit tests will exercise these errors, causing all the logging output as a side-effect, but what really matters to you is whether the unit tests pass or fail.

7 Grading

This project is worth a total of twenty points.

Part 1 (worth 10 points). You must implement several functions and unit tests. Your grader will look to make sure all tests from `SimpleElectionsTestPart1` pass. You get one point for your code passing the Kotlin compiler with no warnings or errors.

Part 2 (worth 10 points). The remaining unit tests, in `SimpleElectionsTestPart2`, are collectively worth 9/10 points. The remaining point is assessed by your human grader, looking at your `README.md` file, where you're asked to answer some questions and build a simple benchmark.